

Benchmarking of Machine Learning and Deep Learning approaches for Neuroimaging Data

Team CogniNet: Antipushina E., Knyshenko M., Kalimullin R., Morozova M.

Mentors: Petr Mokrov, Maria Zubrikhina

Course instructors: Evgeny Burnaev, Alexey Zaytsev

Outline

Theory

- Relevance of the task: problem statement and project goals
- Data description (fMRI, Connectivity matrices and ROIs)

Practical part

- Classical Machine Learning algorithms
- Deep Learning approaches

Relevance of the task

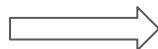
Problem statement:

- At the moment, there is no clear understanding which methods will work best for fMRI data.
- Problem with publication results
- Globally: problem of identifying biomarkers isn't solved.

Project goals:

- Implement classic ML algorithms and fix the metrics.
- Compare with at least one graph approach.

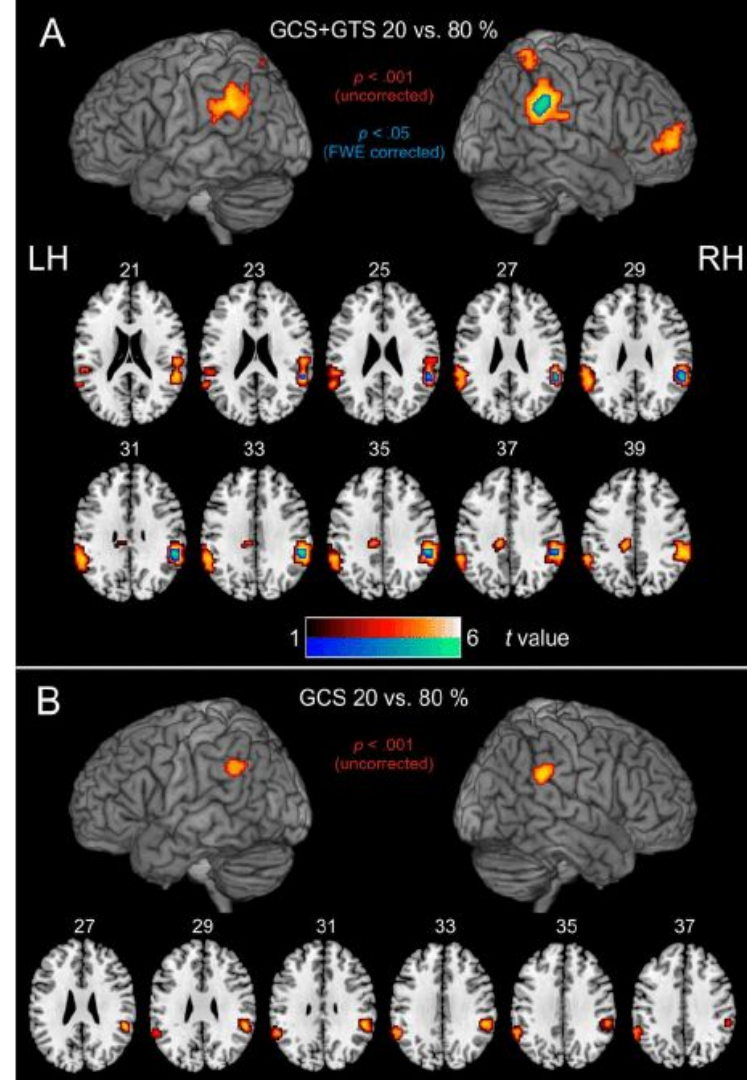
Relevant papers



About fMRI

Functional magnetic resonance imaging (fMRI) measures the small changes in blood flow that occur with brain activity.

It may be used to examine which parts of the brain are handling critical functions, evaluate the effects of stroke or other disease, or to guide brain treatment.



Datasets description

COBRE Dataset+AAL labeling: 163 connectivity matrices (116x116), 79 patients with schizophrenia and 73 healthy controls

ABIDE Dataset+Craddock labeling: 1035 connectivity matrices (200x200), 530 with Autism spectrum disorder (ASD) and 505 healthy controls.

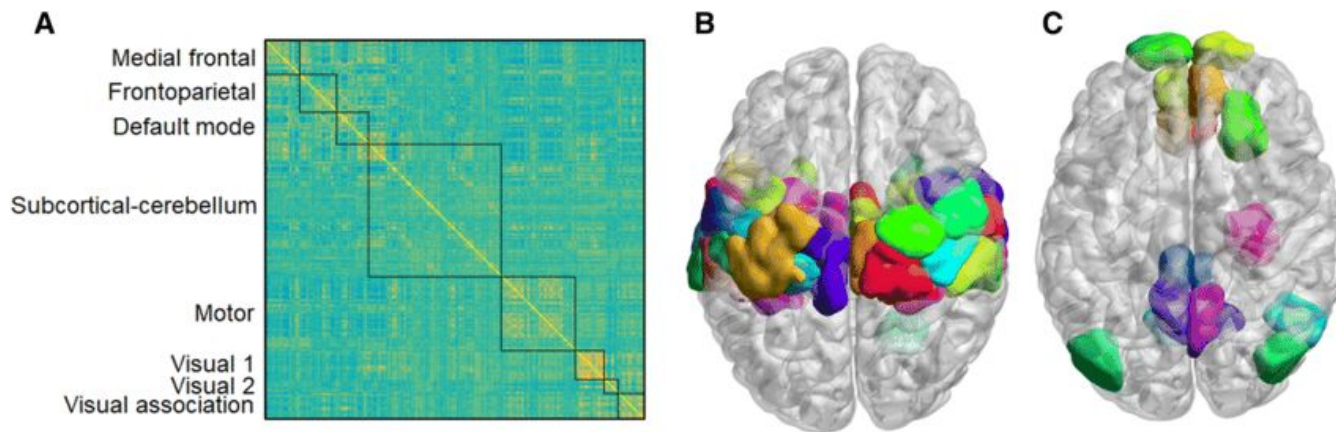


Figure 1 - Brain network-related regions of interest (ROIs)

Classical Machine learning approach

Considered models: Logistic Regression; Random Forest; SVM; SVM + PCA; XGBoost.

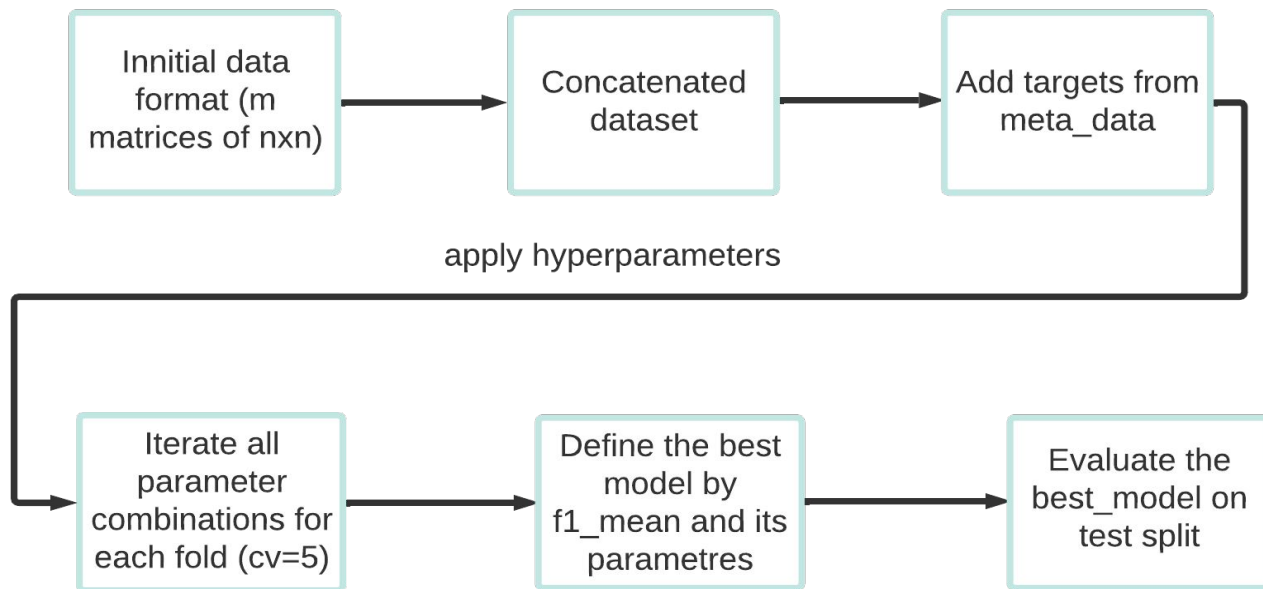


Figure 2 - Code logic block diagram

Classical ML results for COBRE dataset

	F1 score	Accuracy	ROC-AUC
Logistic Regression	0.73 ± 0.056	0.76 ± 0.045	0.76 ± 0.049
Random Forest	0.65 ± 0.073	0.69 ± 0.064	0.69 ± 0.065
SVM	0.76 ± 0.058	0.78 ± 0.039	0.79 ± 0.047
SVM + PCA	0.71 ± 0.099	0.76 ± 0.062	0.75 ± 0.072
XGBoost	0.67 ± 0.044	0.72 ± 0.029	0.71 ± 0.033

Best hyperparameters
'class_weight': None, 'penalty': 'l2', 'C': 0.1, 'solver': 'liblinear'
n_estimators': 50, 'max_depth': 3
'C': 0.01, 'kernel': 'linear', 'probability': True
'C': 10, 'kernel': 'linear', 'class_weight': 'balanced', 'probability': True, n_components=70
'objective': 'binary:logistic', 'eta': 0.1, 'max_depth': 10, 'scale_pos_weight': 1.0, 'alpha': 0

Classical ML results for ABIDE dataset

	F1 score	Accuracy	ROC-AUC
Logistic Regression	0.66 ± 0.013	0.65 ± 0.012	0.66 ± 0.013
Random Forest	0.65 ± 0.037	0.67 ± 0.030	0.67 ± 0.031
SVM	0.67 ± 0.012	0.67 ± 0.013	0.67 ± 0.013
SVM + PCA	0.69 ± 0.017	0.68 ± 0.015	0.68 ± 0.016
XGBoost	0.65 ± 0.007	0.63 ± 0.010	0.63 ± 0.011

Best hyperparameters
'class_weight': None, 'penalty': 'l2', 'C': 10
'n_estimators': 100, 'max_depth': 3
'C': 10, 'kernel': 'rbf'
'C': 0.01, 'kernel': 'linear', 'class_weight': 'balanced', 'probability': True, n_components = 500
'max_depth': 10, 'scale_pos_weight': 1.0, 'alpha': 0

Graph-Attention Neural Network

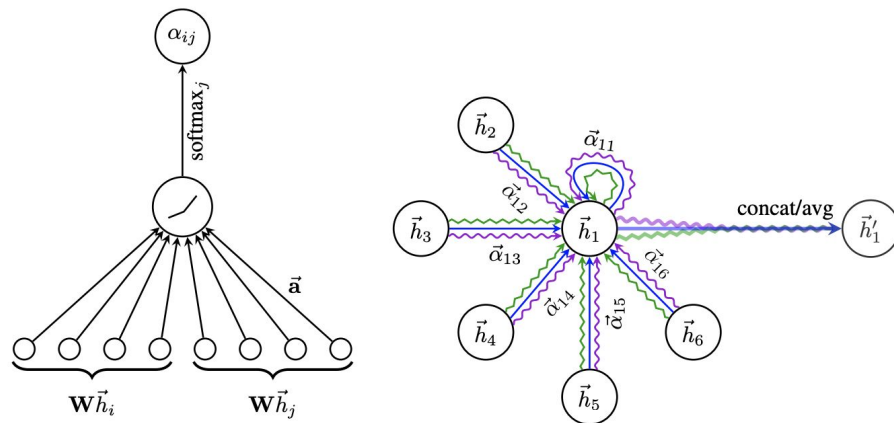


Figure 1: **Left:** The attention mechanism $a(\mathbf{W}\vec{h}_i, \mathbf{W}\vec{h}_j)$ employed by our model, parametrized by a weight vector $\vec{a} \in \mathbb{R}^{2F'}$, applying a LeakyReLU activation. **Right:** An illustration of multi-head attention (with $K = 3$ heads) by node 1 on its neighborhood. Different arrow styles and colors denote independent attention computations. The aggregated features from each head are concatenated or averaged to obtain \vec{h}'_1 .

Figure 3 - GAT architecture

Deep Learning approach. Graph Neural Networks

```
BrainGAT(  
  (convs): ModuleList(  
    (0): Sequential(  
      (0): GATConv(200, 32, heads=2)  
      (1): Linear(in_features=64, out_features=32, bias=True)  
      (2): LeakyReLU(negative_slope=0.2)  
    )  
    (1): Sequential(  
      (0): GATConv(32, 32, heads=2)  
      (1): Linear(in_features=64, out_features=32, bias=True)  
      (2): LeakyReLU(negative_slope=0.2)  
    )  
    (2): Sequential(  
      (0): GATConv(32, 32, heads=2)  
      (1): Linear(in_features=64, out_features=64, bias=True)  
      (2): LeakyReLU(negative_slope=0.2)  
    )  
  )  
  (prepool): Sequential(  
    (0): Linear(in_features=64, out_features=8, bias=True)  
    (1): LeakyReLU(negative_slope=0.2)  
    (2): BatchNorm1d(8, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  )  
  (fcf): BasicMLP(  
    (net): Sequential(  
      (0): Sequential(  
        (0): Linear(in_features=1600, out_features=32, bias=True)  
        (1): LeakyReLU(negative_slope=0.2)  
        (2): Dropout(p=0.6, inplace=True)  
      )  
      (1): Sequential(  
        (0): Linear(in_features=32, out_features=32, bias=True)  
        (1): LeakyReLU(negative_slope=0.2)  
        (2): Dropout(p=0.6, inplace=True)  
      )  
      (2): Sequential(  
        (0): Linear(in_features=32, out_features=2, bias=True)  
      )  
    )  
  )  
)
```



torch_geometric.nn.conv.GATConv
graph attentional operator



fully connected layers (fcf)

number of GATConv layers
number of neurons in GATConv
GATConv parameters (heads)
fcf parameters (number of neurons, layers)
regularizations (BatchNorm, Dropout)

Deep Learning approach. Baseline: multilayer perceptron

```
MLP(  
  (layers): Sequential(  
    (0): Flatten(start_dim=1, end_dim=-1)  
    (1): Linear(in_features=13456, out_features=1024, bias=True)  
    (2): BatchNorm1d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (3): ReLU()  
    (4): Dropout(p=0.1, inplace=False)  
    (5): Linear(in_features=1024, out_features=512, bias=True)  
    (6): ReLU()  
    (7): Dropout(p=0.1, inplace=False)  
    (8): Linear(in_features=512, out_features=2, bias=True)  
  )  
)
```

number of neurons
number of layers
regularizations

Training

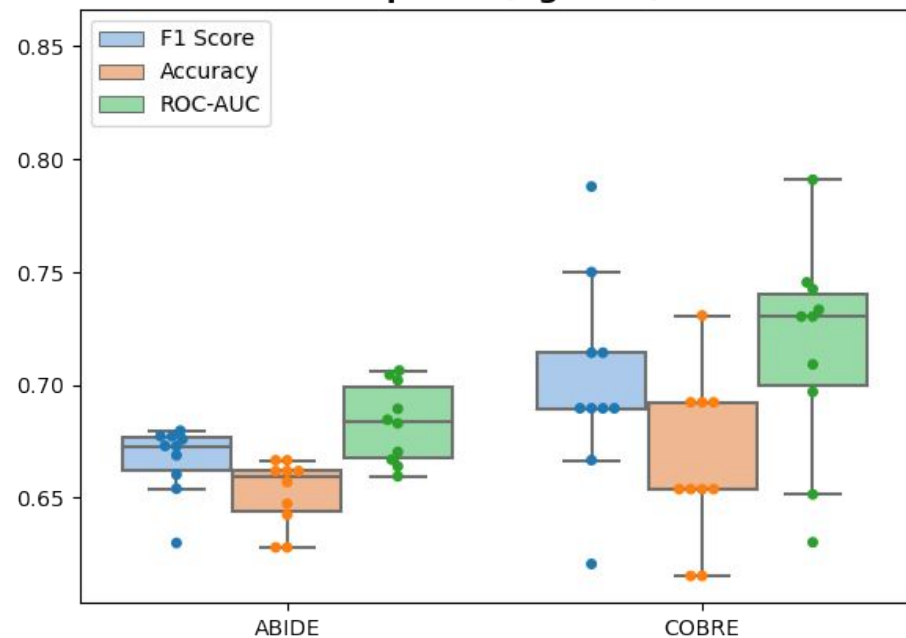
1. Classic **AdamW**, **lr=0.001**, **60 epochs** for training
2. After each epoch, **F1 Score**, **Accuracy**, **ROC-AUC** are estimated on the test set
NB: train test split is the same as in classical ML
3. Best F1 Score, Accuracy, ROC-AUC combination is saved

Repeat NN initialization and training as above **10 times** to estimate variation of metrics

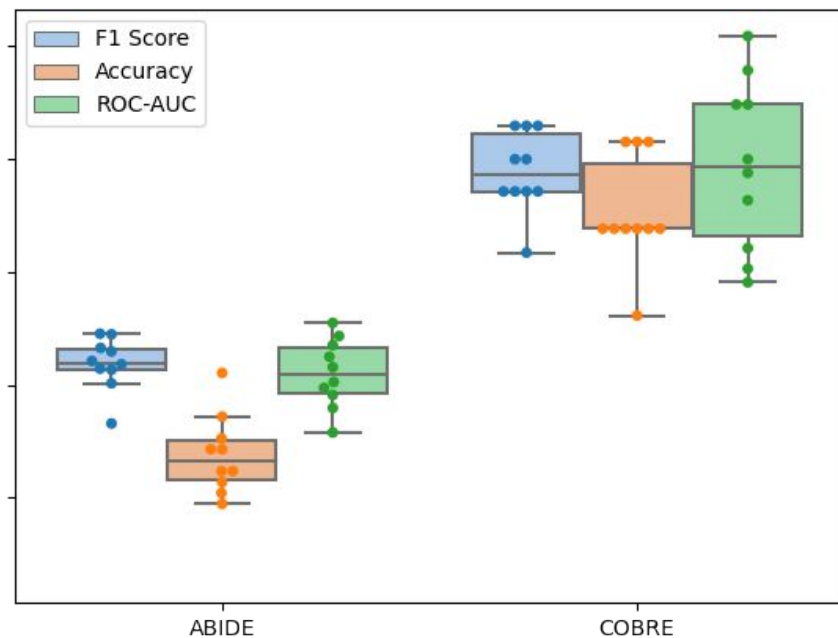
Results

60 epochs, 10 reinitializations
best metric value during training

Graph NN (bgbGAT)



MLP



Comparison of Classic ML and DL

COBRE results

	F1 Score	Accuracy	ROC-AUC
Logistic Regression	0.73 ± 0.056	0.76 ± 0.045	0.76 ± 0.049
Random Forest	0.65 ± 0.073	0.69 ± 0.064	0.69 ± 0.065
SVM	0.76 ± 0.058	0.78 ± 0.039	0.79 ± 0.047
SVM + PCA	0.71 ± 0.087	0.76 ± 0.062	0.75 ± 0.072
XGBoost	0.67 ± 0.044	0.72 ± 0.029	0.71 ± 0.033
bgbGAT	0.70 ± 0.045	0.67 ± 0.036	0.72 ± 0.042
MLP	0.79 ± 0.018	0.78 ± 0.024	0.79 ± 0.038

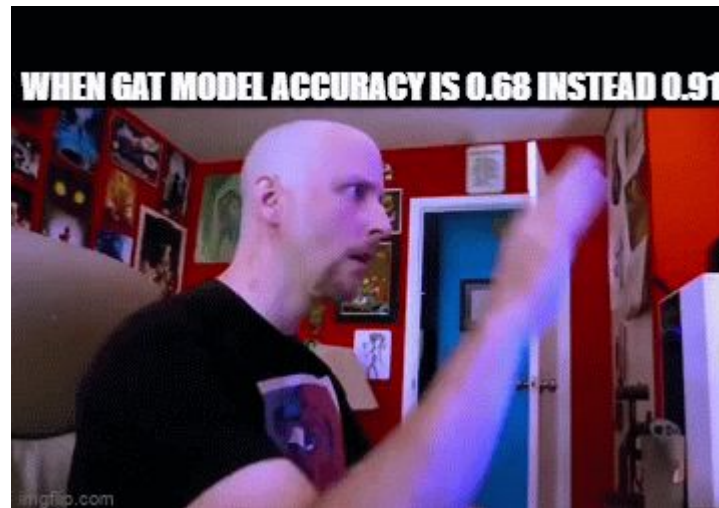
ABIDE results

	F1 Score	Accuracy	ROC-AUC
Logistic Regression	0.66 ± 0.013	0.66 ± 0.013	0.66 ± 0.013
Random Forest	0.65 ± 0.037	0.67 ± 0.030	0.67 ± 0.031
SVM	0.67 ± 0.012	0.67 ± 0.013	0.67 ± 0.013
SVM + PCA	0.69 ± 0.017	0.68 ± 0.015	0.68 ± 0.016
XGBoost	0.65 ± 0.007	0.63 ± 0.010	0.63 ± 0.011
bgbGAT	0.67 ± 0.015	0.65 ± 0.015	0.68 ± 0.017
MLP	0.71 ± 0.012	0.67 ± 0.017	0.71 ± 0.015

Conclusions

1. Usually, we do not have “Big Data” in biomedical research
2. Classic ML approaches are very applicable to biomedical data
3. Well advertised Graph-Attention Neural Networks may not give the best result in biomed
4. For our comparatively small medical datasets, old classic Multilayer Perceptron and SVM beat all the other models

thx!



Link to our github:
<https://github.com/utoprey/CogniNet>



GATConv

The graph attentional operator from the “[Graph Attention Networks](#)” paper —————→



$$\mathbf{x}'_i = \alpha_{i,i} \Theta \mathbf{x}_i + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \Theta \mathbf{x}_j,$$

where the attention coefficients $\alpha_{i,j}$ are computed as

$$\alpha_{i,j} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}^\top [\Theta \mathbf{x}_i \parallel \Theta \mathbf{x}_j]))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}^\top [\Theta \mathbf{x}_i \parallel \Theta \mathbf{x}_k]))}.$$

If the graph has multi-dimensional edge features $\mathbf{e}_{i,j}$, the attention coefficients $\alpha_{i,j}$ are computed as

$$\alpha_{i,j} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}^\top [\Theta \mathbf{x}_i \parallel \Theta \mathbf{x}_j \parallel \Theta_e \mathbf{e}_{i,j}]))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}^\top [\Theta \mathbf{x}_i \parallel \Theta \mathbf{x}_k \parallel \Theta_e \mathbf{e}_{i,k}]))}.$$